

CLIP embedding similarity: within-episode vs across episodes
(intra-episode tightness $\gg$ task gap $\Rightarrow$ selection-invariant retrieval)

